

DS4400 Notes

DS4400 Notes 01/24

1. Convex functions:

A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is convex iff $\forall \theta_1, \theta_2 \in \mathbb{R}^d$ and $\forall \alpha \in [0, 1]$ we have $f(\alpha\theta_1 + (1-\alpha)\theta_2) \leq \alpha f(\theta_1) + (1-\alpha)f(\theta_2)$

In the special case ($d = 1$) $f : \mathbb{R} \rightarrow \mathbb{R}$, f is convex iff $\forall \theta, f''(\theta) \geq 0$

When the function is convex, **local min** \equiv **global min**. When the system is not convex, we might find only a **local min** but not a **global min**

2. Dealing with non convex function:

In gradient descent:

- (a) use larger ρ in the beginning and gradually decrease ρ with iteration.
- (b) Run SGD/GD with multiple random initializations $\theta_1^{(0)}, \theta_2^{(0)} \dots$ and keep the best solution.

3. $\min_{\theta} \sum_{i=1}^N (y_i - \theta^T x_i)^2 \triangleq J(\theta)$

In linear regression, $J(\theta)$ is convex.

4. Robustness of Regression to outliers:

- (a) Run outlier detection algorithm, remove detected outliers, then run Linear Regression on remaining points.
- (b) Robust Regression cost function.
 $\min_{\theta} \sum_{i=1}^N e_i^2$, $e_i \triangleq y_i - \theta^T x_i$
 e^2 is extremely unhappy with large errors.

we might use $|e|$ to replace the function. This might be more tolerance. Then, $\min_{\theta} \sum_{i=1}^N |y_i - \theta^T x_i|$

5. Exercise: $D = \{(x_1, y_1 = 100) \dots (x_1 = 0, y_1 = 0 = 100), (x_{11}, y_{11} = 0), (x_{12}, y_{12} = 0)\}$

$$e^2: 10(\theta - 100)^2 + 2\theta^2 \rightarrow$$

$$\frac{\partial}{\partial \theta} = 20(\theta - 100) + 4\theta = 0 \rightarrow \theta = 83.3$$

$$|e|: \min_{\theta} \sum_{i=1}^{12} |\theta - y_i| = 10|\theta - 100| + 2\theta$$

$$(\theta \leq 100) = \min_{\theta} 10(100 - \theta) + 2\theta$$

$$= 1000 - 8\theta \rightarrow \theta = 100$$

$$(\theta \geq 100) = \min_{\theta} 10(\theta - 100) + 2\theta$$

$$= 12\theta - 1000 \rightarrow \theta = 100$$

6. How to solve l1-norms cost functions?

- (a) No closed form
- (b) we need to be careful with gradient descent
- (c) We need to use convex programming toolboxes (convex optimizations)

7. Huber loss function

$$l_{\delta}(e) = \begin{cases} \frac{1}{2}e^2 & |e| \leq \delta \\ \delta|e| - \frac{\delta^2}{2} & |e| \geq \delta \end{cases}$$

$$\frac{\partial l_{\delta}(e)}{\partial e} = \begin{cases} e & -\delta \leq e \leq \delta \\ \delta & e > \delta \\ -\delta & e < -\delta \end{cases}$$

in huber loss function, we don't have closed form solution but we can run gradient descent now.

8. Definition: Overfitting:

Learning a system from training data that does very well on training data itself (e.g, very low regression error on training data), but performs poorly on test data.

9. **Definition:** Overfitting in Linear Regression

$$\Phi^T \Phi \theta = \Phi^T Y$$

$$\Rightarrow \theta^* = (\Phi^T \Phi)^{-1} \Phi^T Y$$

$$\text{rank}(\Phi^T \Phi) \leq \min\{\text{rk}(\Phi^T), \text{rk}(\Phi)\} = \text{rk}(\Phi) \leq \min\{N, d\}$$

$\Phi^T \Phi$ is $d \times d$ matrix, then rank is $\leq d$.

Therefore, when $N < d$ it is not invertible which means we have multiple solutions and results in overfitting.

DS4400 Notes 01/28

1. **Definition:** Overfitting

Refers to situation where the learned model does well on training data and poorly on testing data.

As d (dimension of system) increases, then training error goes down (can be exactly ZERO for sufficiently large d)

2. In Linear regression:

$$\min \sum_{i=1}^n (\theta^T \phi(x_i) - y_i)^2$$

set the derivative to 0 and we find

$$\Phi^T \Phi \theta = \Phi^T Y$$

Then $\theta^* = (\Phi^T \Phi)^{-1} \Phi^T Y$

When is it the case that $\Phi^T \Phi$ is not invertible?

Since $\Phi^T \Phi \in \mathbb{R}^{N \times d}$

$$\text{rk}(\Phi^T \Phi) \leq \text{rk}(\Phi) \leq \min\{N, d\}$$

$\Phi^T \Phi \in \mathbb{R}^{d \times d}$ is invertible when $\text{rk}(\Phi^T \Phi) = d$. Therefore, when $N < d$, $\text{rk}(\Phi^T \Phi) = N$, $\Phi^T \Phi$ is not invertible. There will be infinitely many solutions for θ .

Generally, need sufficient # samples

3. Test overfitting.

If $\Phi^T \Phi$ is not invertible,

$$\exists v \neq 0, \Phi^T \Phi v = 0$$

$\Rightarrow \theta^* + \alpha v$ is also a solution for any $\alpha \in \mathbb{R}$

$$\begin{aligned} \Phi^T \Phi (\theta^* + \alpha v) &= \Phi^T \Phi \theta^* + \Phi^T \Phi (\alpha v) \\ &= \Phi^T \Phi \theta^* + \alpha \Phi^T \Phi v \\ &= \Phi^T \Phi \theta^* = \Phi^T Y \end{aligned}$$

We can find large α so that θ^* have extremely large entries.

Generally, if the entries are very large (abs) we might have overfitting

4. Treat overfitting

We want to change regression optimization to prevent θ from very large terms.

then we change the cost function:

$$\min_{\theta} \sum_{i=1}^N (\theta^T \phi(x_i) - y_i)^2 + \lambda \sum_{j=1}^d \theta_j^2$$

λ : regularization parameter (> 0)

$\sum_{j=1}^d \theta_j^2$: regularizer.

$\lambda \rightarrow 0$: back to overfitting

$\lambda \rightarrow \infty$: $\theta^* = 0$, underfitting

- (a) closed-form

$$\frac{\partial J}{\partial \theta}$$

$$= 2\Phi^T (\Phi \theta - Y) + \lambda \frac{\partial \sum_{j=1}^d \theta_j^2}{\partial \theta}$$

$$= 2\Phi^T (\Phi \theta - Y) + 2\lambda \theta$$

Let it be zero:

$$\Phi^T \Phi \theta + \lambda \theta = \Phi^T Y$$

$$(\Phi^T \Phi + \lambda I_d) \theta = \Phi^T Y$$

$$\text{Then } \theta^* = (\Phi^T \Phi + \lambda I_d)^{-1} \Phi^T Y$$

- (b) Gradient descent

Find initial $\theta^{(0)}$

$$\theta^t = \theta^{(t-1)} - \rho \frac{\partial J}{\partial \theta} \bigg|_{\theta^{(t-1)}}$$

$$= \theta^{(t-1)} - 2\Phi^T (\Phi \theta^{(t-1)} - Y) + 2\lambda \theta^{(t-1)}$$

5. Hyperparameter Tuning

GD: set learning rate ρ

Robust Reg: Huber loss δ

overfitting and regularization: λ

ρ, δ, λ = hyperparameters

How to pick hyperparameters?

BAD APPROACH 1:

- (a) pick some set of possible $\lambda_i \in \{\lambda_1, \lambda_2 \dots\}$
 Run regression with λ_i and find θ_i^*
 Measure regression error:

$$\epsilon_{tr}(\lambda) = \sum_{i=1}^N ((\theta^*(\lambda))^T x_i - y_i)^2$$

To sum: just find λ for which $\epsilon_{tr}(\lambda)$ is minimum

This approach is setting λ back to 0

Test data needed!!!

- (a) We need to Train λ_i on **training set** to minimize the cost function

$$2\Phi^T(\Phi\theta - Y) + 2\lambda\theta$$

to find θ_i^*

- (b) Measure regression error on the **hold-out set** D^{ho}

$$\epsilon_{tr} = \sum_{x_i, y_i \in D^{ho}} (y_i - (\theta^*(\lambda))^T x_i)^2$$