Neural Networks :

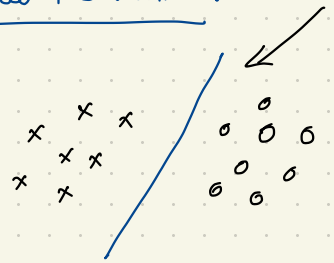

linear function

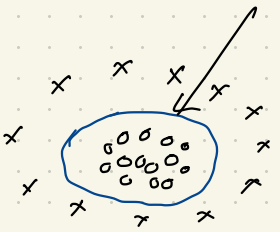

nonlinear function

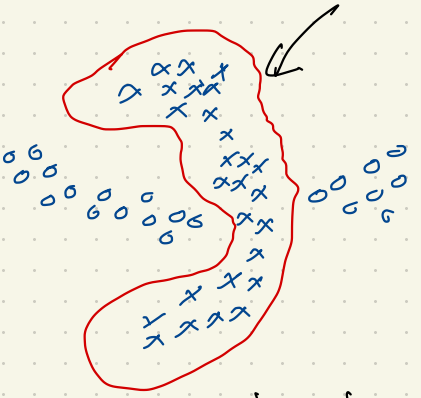

more complex nonlinear function

$$D = \{(x^i, y^i)\}_{i=1}^{N}, \quad x^i \in \mathbb{R}^d, \quad y^i \in \{+1, -1\}$$

find $h_{w,b}(x^i) \simeq y^i, \quad \forall i = 1, \ldots, N$

$$x^i \longrightarrow y^i$$

$$w^T x^i + b \simeq h_{w,b}(x^i)$$

$$w^T \varphi(x^i) + b \simeq h'_{w,b}(x^i)$$

NNs : architectures that allow learning complex functions that can map input data to outputs.

linear / nonlinear

Building blocks of NNs are "neurons"

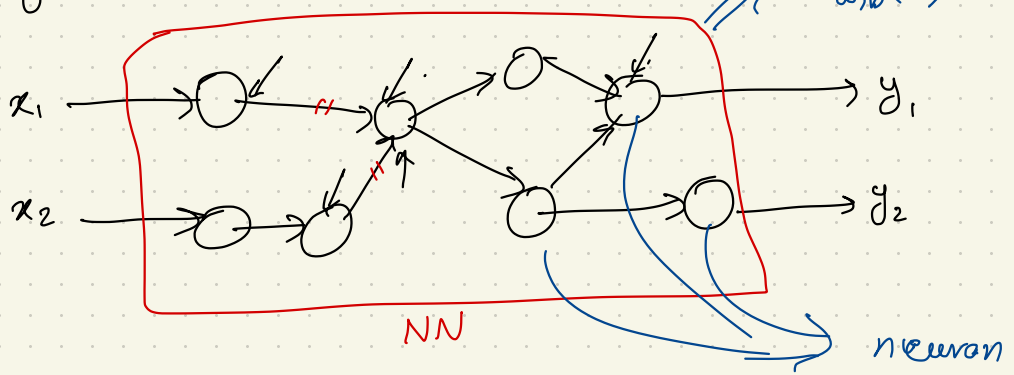

$h_{w,b}(\cdot)$

$\begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \xrightarrow{h_{w,b}(\cdot)} \begin{pmatrix} y_1 \\ y_2 \end{pmatrix}$

A simple neuron can be modeled as follows

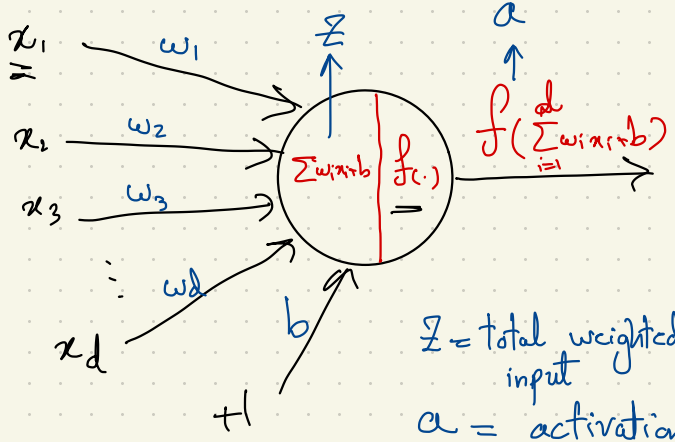

inputs: $x_1, x_2, \cdots, x_d, +1$

parameters: $w_1, w_2, \cdots, w_d, b$

$$\text{output} = f(w_1 x_1 + w_2 x_2 + \cdots + w_d x_d + b)$$

$$= f\left(\sum_{i=1}^{d} w_i x_i + b\right)$$

z = total weighted input

a = activation / output

Common choices of activation functions :

* Sigmoid function $f(z) = \sigma(z) = \frac{1}{1+e^{-z}}$

$\in [0,1]$

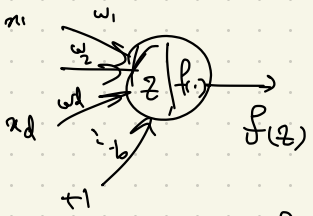

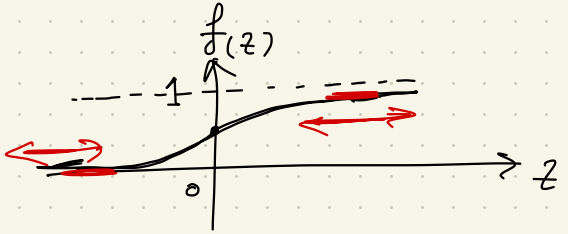

$z \to \infty \Rightarrow +1$

$z = 0 \Rightarrow 0.5$

$z \to -\infty \Rightarrow 0$

Good for modeling binary outputs (spam or not)

probabilities (p(email is spam) , p('car in image')

scores that are bounded (student grade [0,100] ↓ [0, 1])

* Tangent Hyperbolic function

$$f(z) = \frac{e^{z} - e^{-z}}{e^{z} + e^{-z}}$$

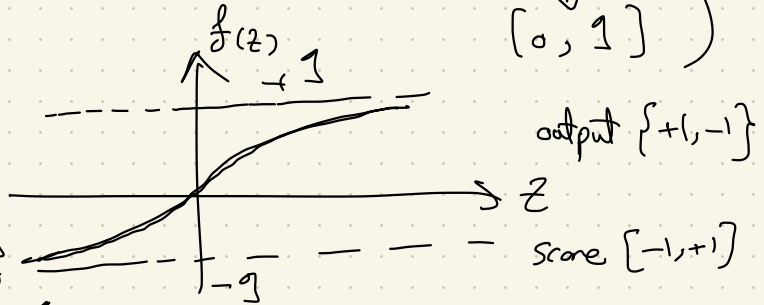

$z \to \infty \Rightarrow 1$

$z = 0 \Rightarrow 0$

$z \to -\infty \Rightarrow -1$

* Rectifier Linear activation function (ReLu)

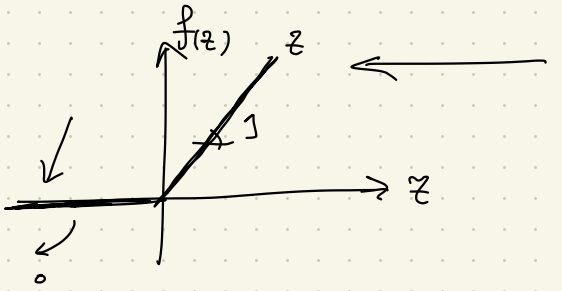

$$Relu(z) = \max\{0, z\}$$

$$= \begin{cases} z & z \geq 0 \\ 0 & z \leq 0 \end{cases}$$

$$\Rightarrow \frac{\partial Relu(z)}{\partial z} = \begin{cases} 1 & z > 0 \\ \text{undefined} & z = 0 \\ 0 & z < 0 \end{cases}$$

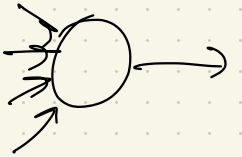

$\Rightarrow$ good for modeling outputs $[0, \infty)$ $\rightarrow$ stock value $

temperature (Kelvin)

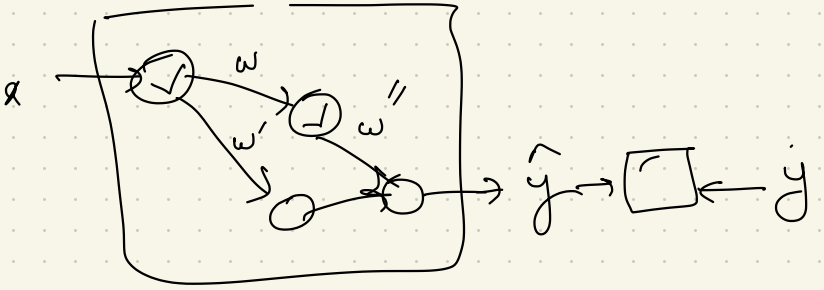

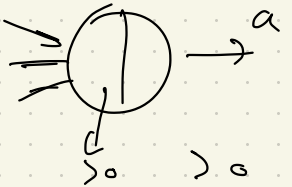

* Leaky Relu function

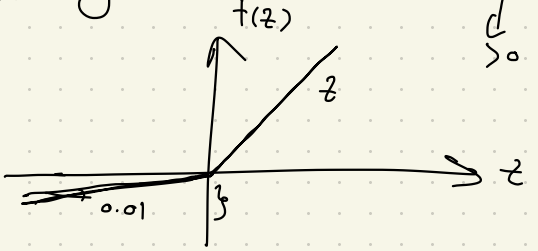

Relu is more biologically plausible!

Biological Neuron

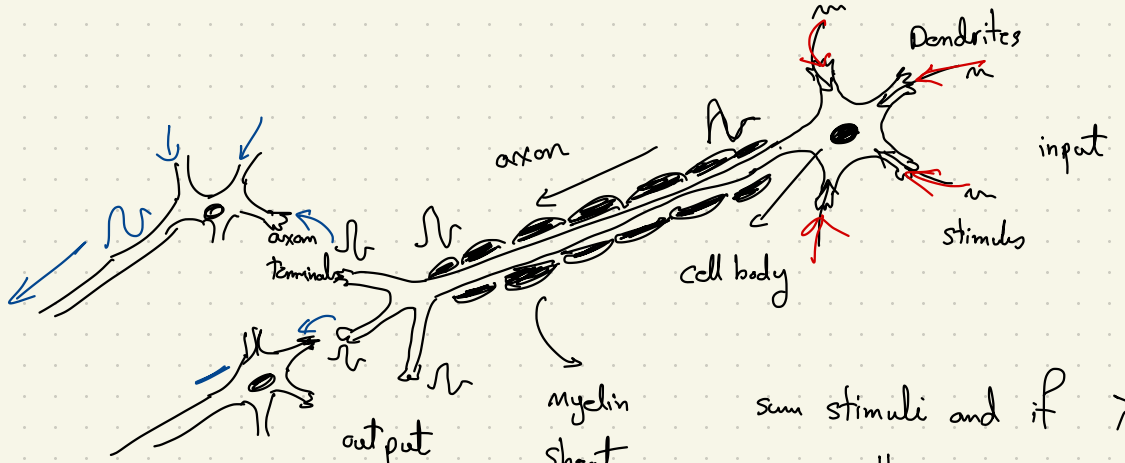

sum stimuli and if > thr

⇓

fire

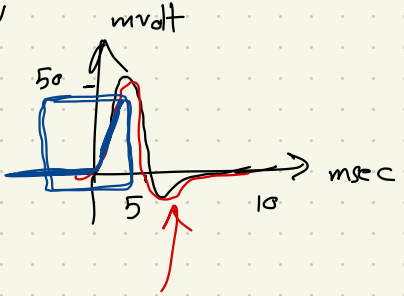

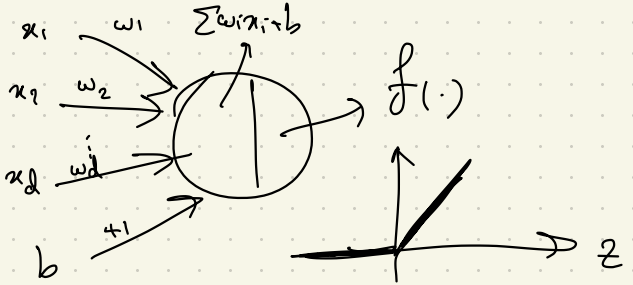

Back to Artificial Neural Networks :

Feedforward NNs :

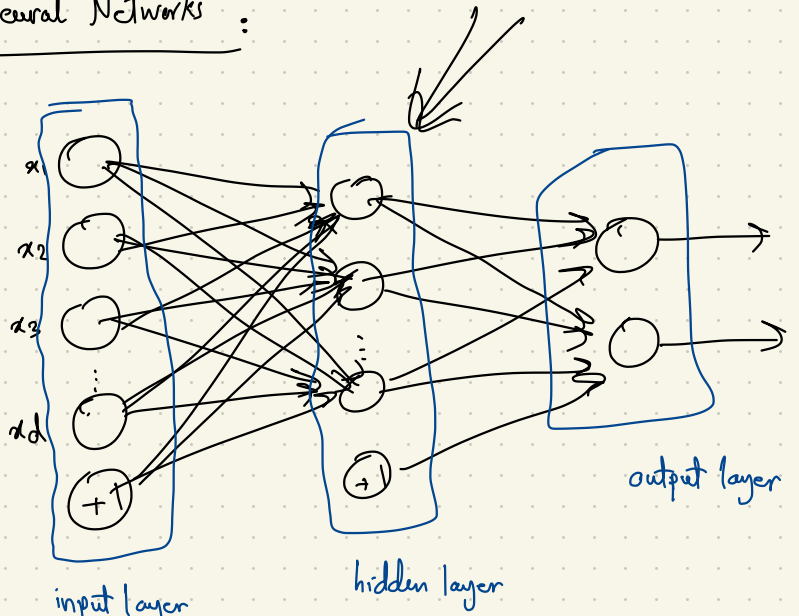

Shallow

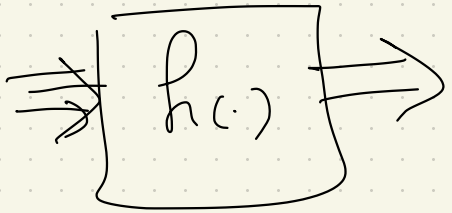

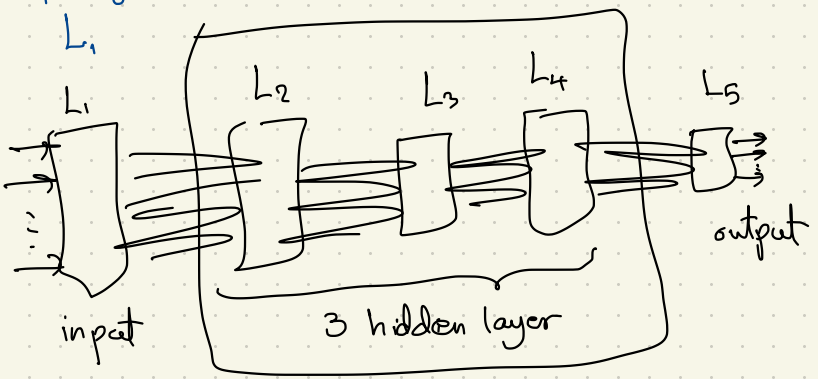

Deep

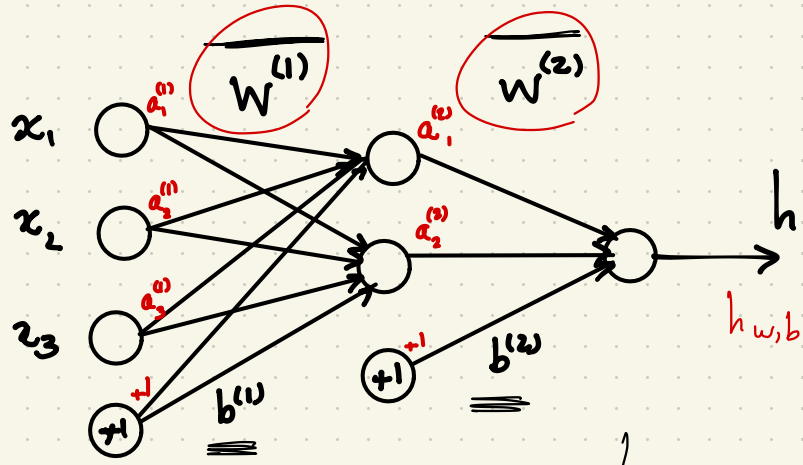

x^{test}

$D = \{(x^i, y^i)\}_{i=1}^{N}$

$h_{w,b}(x^i) \approx y^i$ car/not car

$h_{w,b}(x^{test})$

x^i ⇐ Image

L_1 $S_1 = 3$

L_2 $S_2 = 2$

L_3 $S_3 = 1$

n = # layers

S_l = # nodes in layer l (excluding intercept/+1)

Want to learn all weights and biases between different layers using training data!

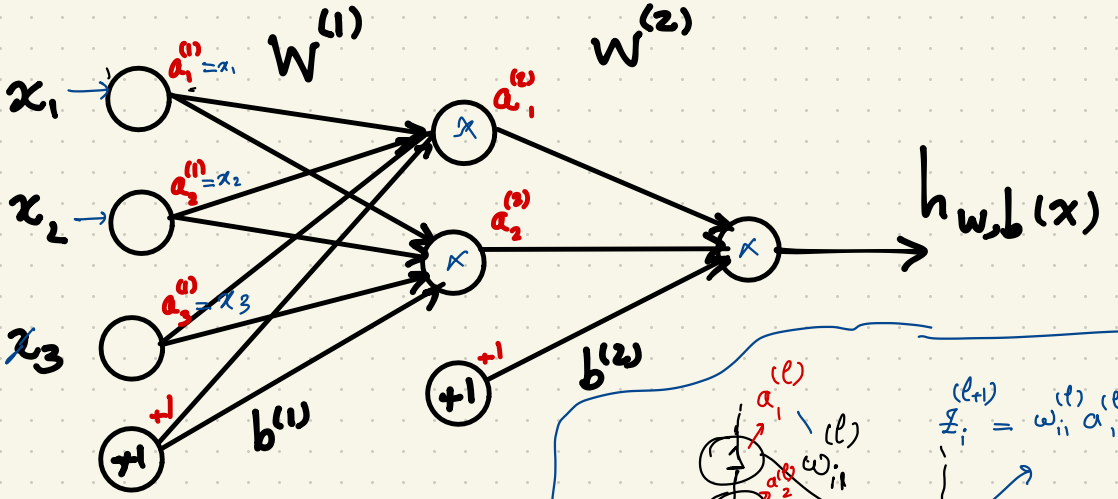

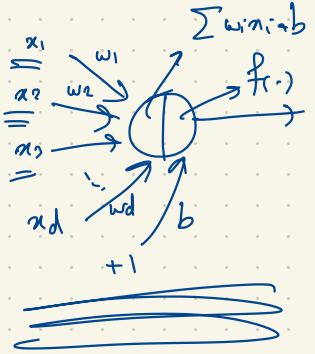

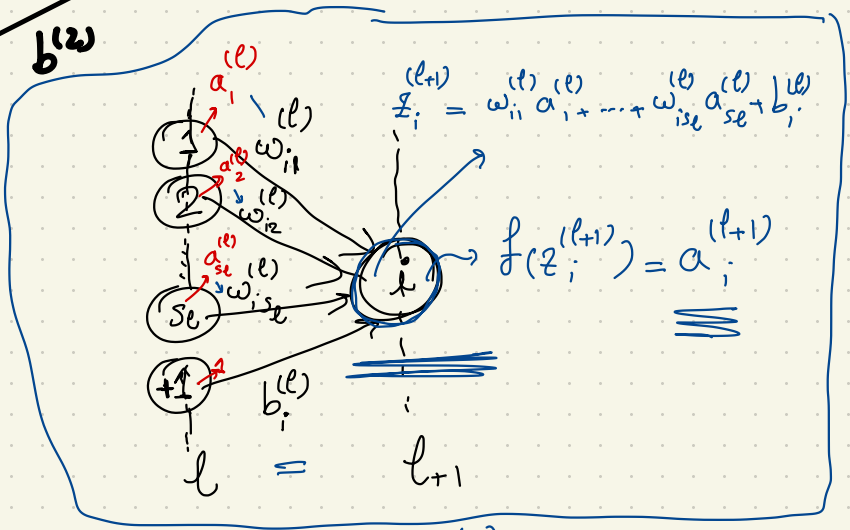

In layer 1: $z_i^{(1)} = a_i^{(1)} = x_i$